

1 Background

Comprehensive genomic profiling (CGP) guides treatment in NSCLC, but panel content, cost and turnaround time differ widely between institutions, and smaller targeted panels leave part of the tumour genome unmeasured.

EBAI class B. Under the ESMO basic requirements for AI-based biomarkers (EBAI), a model that predicts an established biomarker from an input different from the gold-standard test is class B; here the input is the smaller panel itself. Because the input differs from the reference test, risk and complexity exceed class A. Such tools may predict only a subset of alterations with sufficient sensitivity and specificity, and usually aim to **enrich** a population for confirmatory testing, or, where a negative result reliably excludes the alteration, to **rule out** testing. They do not substitute an established genomic test.

Objective. To learn the conditional distribution of alterations in genes absent from a targeted panel given the genes it measures, benchmark it against simple baselines, and evaluate it along the seven EBAI dimensions: comparator, performance, generalisability, fairness, explainability, cost and turnaround time.

IMPACT341 · 341 genes (input)

164 genes (target)

Figure 1. Task definition. IMPACT341 ⊂ IMPACT505; the 164 genes present only in IMPACT505 are masked from the input and predicted.

Table 1. Cohort and target characteristics (MSK-CHORD).	
Patients, NSCLC, sequenced on IMPACT505	2,226
Input genes (IMPACT341)	341
Target genes (IMPACT505 \ IMPACT341)	164
Altered target gene–patient pairs	1.21%
Test-set patients	≈334
Patient-level split: train / validation / test	70 / 15 / 15%*

* Test (15%, ≈334 patients) and validation (15%) as in the abstract; the remaining 70% trains the model. The abstract’s “75%” for training would sum to 105%.

Alteration encoding. Target: one binary event per gene and patient (any non-synonymous mutation, including variants of unknown significance, amplification or deep deletion; fusions and other structural variants not modelled). Input: separate mutation, amplification and deletion indicators for the 341 assayed genes. One sample per patient. Genes not covered by a sample’s panel are treated as missing, never as wild type.

3 Methods

3.1 Problem formulation. Given the observed IMPACT341 profile c , we learn the conditional distribution $p(x_1 \mid c)$ over the 164 target genes. Binary targets are embedded as $x_1 \in \{-1, +1\}^{164}$ (continuous relaxation); generated profiles are binarised by sign. Data were split by patient into 70% training, 15% validation (all model and hyper-parameter selection) and 15% test, evaluated once.

3.2 Conditional flow matching. A velocity field v_θ , a neural network (architecture variants compared on the validation set), is regressed onto the straight-line path between Gaussian noise and data:

$$x_t = (1 - t) x_0 + t x_1, \quad x_0 \sim \mathcal{N}(0, I), \quad t \sim \mathcal{U}(0, 1)$$
$$\mathcal{L}(\theta) = \mathbb{E} \parallel v_\theta(x_t, t \mid c) - (x_1 - x_0) \parallel^2$$

$x_0 \sim \mathcal{N}(0, I)$

$\rightarrow$

$dx/dt = v_\theta(x_t, t \mid c)$

$\rightarrow$

$\hat{x}_1 \in \mathbb{R}^{164}$

conditioning c : observed IMPACT341 profile

Figure 2. Sampling: the learned ODE, integrated numerically from $t = 0$ to 1, transports noise to a completed 164-gene profile conditioned on the observed panel.

3.3 Probability read-out. At $t = 0$, x_t is independent of x_1 , so $v^*(x_0, 0 \mid c) = \mathbb{E}[x_1 \mid c] - x_0$ and $\hat{p} = \frac{1}{2}(1 + x_0 + v\theta(x_0, 0 \mid c))$, averaged over K noise draws. This avoids the Monte Carlo ties that sample frequencies produce at ~1% prevalence.

3.4 Evaluation (EBAI). Pre-specified primary endpoint: pooled AUROC. Secondary: per-gene AUROC (genes with ≥ 5 test events; prevalence-only = 0.50); AUPRC (chance = 0.012); calibration plot, slope, intercept, O:E and Brier skill; at the intended-use threshold for ≥ 1 altered target gene, sensitivity, specificity, PPV, NPV and LR±; NRI over the burden + locus model and decision-curve analysis; joint fidelity (energy score versus independent Bernoulli draws). Comparators (same splits): prevalence, burden + locus, L2 logistic regression; 95% CIs and paired differences by patient bootstrap.

Sample size (post hoc; EBAI asks a priori): ~334 test patients give a 95% CI half-width of $\approx \pm 0.05$ for a patient-level AUROC of 0.79 (Hanley–McNeil, 43% prevalence assumed).

3.5 External replicate. MSK-IMPACT 2017 (cBioPortal): IMPACT341 → 69 IMPACT410-only genes, 1,215 NSCLC patients, 3 × 5-fold patient-level CV, AUROC within folds.

4 Results

4.1 Primary cohort. Across architectures and training paradigms, the flow model reached a pooled test AUROC of 0.77–0.79, above the prevalence-only baseline in every variant. The abstract reports no other metric; Table 2 sets these gaps against what the replicate measures.

Metric (EBAI)	MSK-CHORD	Replicate
Pooled AUROC	0.77–0.79	0.778 (0.754–0.796)
Per-gene AUROC	<i>n.r.</i>	0.728 (0.705–0.753)
AUPRC (chance $\approx$ 0.011)	<i>n.r.</i>	0.114
Calibration slope · O:E	<i>n.r.</i>	0.97 · 1.08
Brier skill vs prevalence	<i>n.r.</i>	0.042
NPV at 90% sensitivity*	<i>n.r.</i>	0.86 (LR– 0.22)
Joint fidelity (energy score)	<i>n.r.</i>	not applicable

* Patient level: ≥ 1 altered target gene (prevalence 43%).

4.2 Benchmark on public data. Pooled AUROC compares pairs across genes, so ranking genes by training frequency alone scores 0.715. Per gene, a 7-feature model matches 1,026-feature logistic regression (0.760 vs 0.728), and pan-cancer training adds +0.039 (95% CI 0.024–0.057; Figure 3).

Category	Pooled AUROC	Per-gene AUROC
Prevalence only	0.715	0.500
Burden + locus	0.760	0.760
L2 logistic regression	0.778	0.728
Pan-cancer logistic	0.808	0.767

Figure 3. Replicate ($n = 1,215$): AUROC with patient-bootstrap 95% CI. Pan-cancer: trained on 7,773 patients with cancer type as covariate.

4.3 Simulated versus real input. In 167 patients sequenced on both panels, taking the input from the separate IMPACT341 assay instead of masking the IMPACT410 sample lowered pooled AUROC from 0.84 to 0.70 and per-gene AUROC from 0.80 to 0.61 (pan-cancer logistic regression). The paired samples also differ in specimen and time, so the gap combines assay shift and tumour heterogeneity.

4.4 Sources of signal. Amplifications were highly discriminable (AUROC 0.95), explained largely by co-amplification with neighbouring assayed genes (NFKBIA with NKX2-1 in 39/39 cases). Mutations reached 0.73, no better than three mutational-burden counts alone (Figure 4).

Event Type	Per-gene AUROC
Amplifications (134 events)	0.95
Mutations (646 events)	0.73
Mutations, burden only	0.73
Deep deletions (18 events)	0.71

Figure 4. Replicate: per-gene AUROC by event type (positives-weighted mean; axis origin 0.50), pan-cancer logistic regression.

5 EBAI assessment

EBAI class B · indirect measure of existing biomarkers · use: enrichment			
Essential	Comparator	Met	IMPACT505 NGS on the same sample
	Discrimination	Partial	Pooled AUROC 0.77–0.79; per gene n.r.
	Operating point	Partial	Replicate: Se 0.90, NPV 0.86, LR– 0.22
	Calibration	Partial	Replicate: slope 0.97, O:E 1.08
	Indep. cohort	Not yet	Single institution; external cohort needed
	Assay variation	Partial	Real IMPACT341 input: AUROC 0.84 → 0.70
Recommended	Fairness	Not yet	Sex, ancestry, smoking, histology
	Explainability	Met	Ablation: CN locus + mutational burden
Additional	Cost	Partial	Reuses routine panel data; not costed
	Turnaround	Partial	Inference in seconds; workflow not timed
Validation	Sample size	Partial	Post hoc: $\approx \pm 0.05$ AUROC with 334 test patients
	NRI / DCA	Not yet	Versus burden + locus model; triage threshold
Post-adoption	Monitoring	Not yet	Drift with panel versions and populations

Figure 5. Assessment along the seven EBAI dimensions and validation requirements for class B. Replicate values from MSK-IMPACT 2017; n.r., not reported. CN, copy number; DCA, decision-curve analysis; NRI, net reclassification improvement.

6 Limitations

Abstract

Evaluation. Single 15% test split; pooled AUROC without CI, per-gene AUROC, AUPRC or calibration; split stated as 75/15/15%.

Validity. Simulated masking of one IMPACT505 assay; single institution, no independent or real paired-assay cohort.

Scope. Binary per-gene events ignore variant level (e.g. KRAS G12C) and fusions; actionable genes and subgroups not analysed.

Generative claim. No joint-fidelity metric; architecture variants unspecified.

Poster and replicate

Replicate. 69 genes, 2014–2016 MSK data: not independent of MSK-CHORD under EBAI.

Models. Operating point and calibration from L2 logistic regression, not the flow model; sample size post hoc.

7 Conclusions

- A targeted panel predicts part of the unassayed genome, mainly through copy-number co-location and mutational burden.
- As an EBAI class B biomarker its use is enrichment for confirmatory CGP, not rule-out (replicate NPV 0.86).
- Next: per-gene and calibration metrics, an independent multi-centre cohort and a fairness audit.

References

1. Jee J, et al. Automated real-world data integration improves cancer outcome prediction. Nature 2024;636:728–36.
2. Zehir A, et al. Mutational landscape of metastatic cancer revealed from prospective clinical sequencing of 10,000 patients. Nat Med 2017;23:703–13.
3. Lipman Y, et al. Flow matching for generative modeling. ICLR 2023. 4. Collins GS, et al. TRIPOD+AI statement. BMJ 2024;385:e078378.
5. Aldea M, et al. ESMO basic requirements for AI-based biomarkers in oncology (EBAI). Ann Oncol 2026;37:414–25.

Data and code

MSK-CHORD and msk_impact_2017 via cBioPortal. Code and split IDs: github.com/chiaralouisa/brainmets-POP, available on request.

Disclosures · Contact

The authors declare no conflicts of interest.
Louisa Hempel · louisa.hempel@googlemail.com